



CS 559: NEURAL NETWORKS

---

## Homework 5

---

*Author:*

Gaetano Coppoletta

*Email:*

[gcoppo2@uic.edu](mailto:gcoppo2@uic.edu)

October 20, 2022

## 1 Introduction

In this homework we have to design and train a neural network to classify images containing various shape. The dataset contains 90,000 images, where each image has size 200x200 and belongs to one of the 9 classes: Circle, Square, Octagon, Heptagon, Nonagon, Star, Hexagon, Pentagon, Triangle. There are 10,000 images per class.

## 2 Create training and testing sets

First of all we want to write a program that reads the files and creates one variable for training and another for testing. The training set should contain 8,000 images per class, and the test set should contain the remaining 2,000 images per class, with the corresponding labels as indicated by the file names. We use the file names to split training and testing sets.

## 3 Design the neural network

Now we want to design and implement a neural network that will take a 200x200 image and decide which one of the 9 classes the input corresponds to.

### 3.1 First network

First of all we try the neural network implemented in torch3.py, adapting some parameters and training it for 9 epochs. This network uses Adam optimizer as optimization method, training batch size of 100, test batch size equal to 1000 and Cross Entropy as loss function. The code of the network is listed below.

```
1 class Net(nn.Module):
2     def __init__(self):
3         super(Net, self).__init__()
4         self.conv1 = nn.Conv2d(3, 32, 3, 1)
5         self.conv2 = nn.Conv2d(32, 64, 3, 1)
6         self.dropout1 = nn.Dropout(0.25)
7         self.dropout2 = nn.Dropout(0.5)
8         self.fc1 = nn.Linear(614656, 128)
9         self.fc2 = nn.Linear(128, 10)
10
11     def forward(self, x):
12         x = self.conv1(x)
13         x = F.relu(x)
14         x = self.conv2(x)
15         x = F.relu(x)
16         x = F.max_pool2d(x, 2)
17         x = self.dropout1(x)
18         x = torch.flatten(x, 1)
19         x = self.fc1(x)
20         x = F.relu(x)
21         x = self.dropout2(x)
22         x = self.fc2(x)
```

23 `return x`

The results are listed below.




Figure 1: Accuracy and loss on train and test sets

As we can see from the chart, this network is not working as desired, in fact we have a test accuracy which is less than 40

### 3.2 Second network

After doing some tests, in particular adding, removing and modifying the layers of the network, a good solution was found using three convolutional layer, three max pooling layer and 3 linear layer, using relu as activation function. This network uses Adam optimizer as optimization method, training batch size of 100, test batch size equal to 100 and Cross Entropy as loss function. This network has been trained for 20 epochs and it uses a learning parameter equal to  $10^{-3}$ . With this network we obtain better results, at the end of epoch 20 we have:

- Training Loss: 0.373081, Training Accuracy: 85.81
- Test Loss: 0.451795, Test Accuracy: 83.95

The code of the network is listed below:

```
1 class Net(nn.Module):
2     def __init__(self):
3         super(Net, self).__init__()
4         self.conv1 = nn.Conv2d(3, 8, 5)
5         self.max_pool1 = nn.MaxPool2d(4, 4)
6         self.conv2 = nn.Conv2d(8, 16, 5)
7         self.max_pool2 = nn.MaxPool2d(5, 5)
8         self.conv3 = nn.Conv2d(16, 32, 5)
9         self.max_pool3 = nn.MaxPool2d(5, 5)
10        self.lin1 = nn.Linear(32, 128)
11        self.lin2 = nn.Linear(128, 84)
12        self.lin3 = nn.Linear(84, 9)
13    def forward(self, x):
14        x = self.conv1(x)
15        x = F.relu(x)
16        x = self.max_pool1(x)
17        x = self.conv2(x)
18        x = F.relu(x)
19        x = self.max_pool2(x)
20        x = self.conv3(x)
21        x = F.relu(x)
22        x = self.max_pool3(x)
23        x = torch.flatten(x, 1)
24        x = self.lin1(x)
25        x = F.relu(x)
26        x = self.lin2(x)
27        x = F.relu(x)
28        x = self.lin3(x)
29        return x
```

The results are listed below.




Figure 2: Accuracy and loss on train and test sets

## 4 Code 1

```
1 import argparse
2 import torch
3 import torch.nn as nn
4 import torch.nn.functional as F
5 import torch.optim as optim
6 import torchvision
7 import matplotlib.pyplot as plt
8 from torchvision import datasets, transforms
9 from torch.optim.lr_scheduler import StepLR
10 import os
11 import shutil
12
13 from torchvision.datasets import ImageFolder
14
15 names = ["Circle", "Square", "Octagon", "Heptagon", "Nonagon", "Star", "Hexagon",
16         , "Pentagon", "Tringle"]
17 train_loss=[]
18 train_acc=[]
19 test_loss=[]
20 test_acc=[]
21
22 class Net(nn.Module):
23     def __init__(self):
24         super(Net, self).__init__()
25         self.conv1 = nn.Conv2d(3, 8, 5)
26         self.max_pool1 = nn.MaxPool2d(4, 4)
27         self.conv2 = nn.Conv2d(8, 16, 5)
28         self.max_pool2 = nn.MaxPool2d(5, 5)
29         self.conv3 = nn.Conv2d(16, 32, 5)
30         self.max_pool3 = nn.MaxPool2d(5, 5)
31         self.lin1 = nn.Linear(32, 128)
32         self.lin2 = nn.Linear(128, 84)
33         self.lin3 = nn.Linear(84, 9)
34     def forward(self, x):
```

```
34         x = self.conv1(x)
35         x = F.relu(x)
36         x = self.max_pool1(x)
37         x = self.conv2(x)
38         x = F.relu(x)
39         x = self.max_pool2(x)
40         x = self.conv3(x)
41         x = F.relu(x)
42         x = self.max_pool3(x)
43         x = torch.flatten(x, 1)
44         x = self.lin1(x)
45         x = F.relu(x)
46         x = self.lin2(x)
47         x = F.relu(x)
48         x = self.lin3(x)
49         return x
50 #create directories
51 os.mkdir("images")
52 os.mkdir("./images/test")
53 os.mkdir("./images/train")
54
55 for i in range (0,len(names)):
56     os.mkdir("./images/test/"+names[i])
57     os.mkdir("./images/train/"+names[i])
58
59 i=0
60 j=0
61 names.sort()
62 filenames = os.listdir("./output") #where the images are stores
63 filenames.sort()
64 for filename in filenames:
65     if i < 8000:
66         shutil.copy("./output/"+filename, "./images/train/"+names[j]+"/")
67         i+=1
68     elif i<10000:
69         shutil.copy("./output/"+filename, "./images/test/"+names[j]+"/")
70         i+=1
71     else:
72         j+=1
73         i=0
74         if j==9:
75             break
76
77
78
79 def train(args, model, device, train_loader, optimizer, epoch):
80     model.train()
81     tot_loss = 0
82     correct = 0
83     for batch_idx, (data, target) in enumerate(train_loader):
84         data, target = data.to(device), target.to(device)
85         optimizer.zero_grad()
86         output = model(data)
```

```
87     loss = torch.nn.CrossEntropyLoss()(output, target)
88     loss.backward()
89     optimizer.step()
90
91     pred = output.argmax(dim=1, keepdim=True)
92     correct += pred.eq(target.view_as(pred)).sum().item()
93
94     tot_loss = tot_loss + loss.item()
95     if batch_idx % args.log_interval == 0:
96         print('Train Epoch: {} [{}/{}] ({:.0f}%) \tLoss: {:.6f}, Accuracy:
97               {:.2f}%'.format(
98             epoch, batch_idx * len(data), len(train_loader.dataset),
99             100. * batch_idx / len(train_loader), tot_loss / (
100                 batch_idx + 1),
101             100.0 * correct / ((batch_idx + 1) * args.batch_size)))
102
103     print('End of Epoch: {}'.format(epoch))
104     train_loss.append(tot_loss / (len(train_loader)))
105     train_acc.append(100.0 * correct / (len(train_loader) * args.batch_size))
106     print('Training Loss: {:.6f}, Training Accuracy: {:.2f}%'.format(
107         tot_loss / (len(train_loader)), 100.0 * correct / (len(train_loader) *
108             args.batch_size)))
109
110 def test(args, model, device, test_loader):
111     model.eval()
112     tot_loss = 0
113     correct = 0
114     with torch.no_grad():
115         for data, target in test_loader:
116             data, target = data.to(device), target.to(device)
117             output = model(data)
118             tot_loss += torch.nn.CrossEntropyLoss()(output, target).item() #
119                 sum up batch loss
120             pred = output.argmax(dim=1, keepdim=True) # get the index of the
121                 max log-probability
122             correct += pred.eq(target.view_as(pred)).sum().item()
123     test_loss.append(tot_loss / (len(test_loader)))
124     test_acc.append(100.0 * correct / (len(test_loader) * args.test_batch_size))
125     print('Test Loss: {:.6f}, Test Accuracy: {:.2f}%'.format(
126         tot_loss / (len(test_loader)), 100.0 * correct / (len(test_loader) *
127             args.test_batch_size)))
128
129 def main():
130     # Training settings
131
132     parser = argparse.ArgumentParser(description='HW5')
133     parser.add_argument('--batch-size', type=int, default=100, help='input batch
134         size for training (default: 100)')
135     parser.add_argument('--test-batch-size', type=int, default=100,
136         help='input batch size for testing (default: 100)')
137     parser.add_argument('--epochs', type=int, default=20, help='number of epochs
```

```
        to train')
133 parser.add_argument('--lr', type=float, default=1e-3, help='learning rate')
134 parser.add_argument('--gamma', type=float, default=0.7, help='Learning rate
    step gamma')
135 parser.add_argument('--seed', type=int, default=2022)
136 parser.add_argument('--log-interval', type=int, default=100,
137                     help='how many batches to wait before logging training
                            status')
138 parser.add_argument('--save-model', action='store_true', default=True, help=
    'For Saving the current Model')
139 parser.add_argument('-f')
140 args = parser.parse_args()
141
142 torch.manual_seed(args.seed)
143
144 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
145
146 transform = transforms.Compose([transforms.ToTensor()])
147
148 dataset1 = datasets.ImageFolder('./images/train',transform)
149 dataset2 = datasets.ImageFolder('./images/test',transform)
150 train_loader = torch.utils.data.DataLoader(dataset1, batch_size=100, shuffle
    =True)
151 test_loader = torch.utils.data.DataLoader(dataset2, batch_size=100)
152
153 model = Net().to(device)
154 optimizer = optim.Adam(model.parameters(), lr=args.lr)
155
156 scheduler = StepLR(optimizer, step_size=1, gamma=args.gamma)
157 for epoch in range(1, args.epochs+1):
158     train(args, model, device, train_loader, optimizer, epoch)
159     test(args, model, device, test_loader)
160     scheduler.step()
161
162 if args.save_model:
163     torch.save(model.state_dict(), "/content/drive/MyDrive/Colab Notebooks
        /0602-657811153-Coppoletta.pt")
164
165 plt.plot(train_loss, c = 'b')
166 plt.plot(test_loss, c = 'r')
167 plt.legend(['Train Loss', 'Test Loss'])
168 plt.title("Train Loss & Test Loss")
169 plt.xlabel("Epochs")
170 plt.ylabel("Loss")
171 plt.savefig('/content/drive/MyDrive/Colab Notebooks/loss', format = 'eps')
172 plt.show()
173 plt.plot(train_acc, c = 'b')
174 plt.plot(test_acc, c = 'r')
175 plt.legend(['Train Accuracy', 'Test Accuracy'])
176 plt.title("Train Accuracy & Test Accuracy")
177 plt.xlabel("Epochs")
178 plt.ylabel("Accuracy")
179 plt.savefig('/content/drive/MyDrive/Colab Notebooks/accuracy', format = 'eps')
```



```
        ')
180     plt.show()
181
182
183
184 if __name__ == '__main__':
185     main()
```

## 5 Code 2

```
1 import os
2 import torch
3 import torchvision
4 import torchvision.transforms as transforms
5 import torch.nn as nn
6 import torch.nn.functional as F
7 from PIL import Image
8
9 name = []
10 image = []
11 image_path = './images' #current directory, where the images are
12 for i in os.listdir(image_path):
13     name.append(i)
14     image.append(Image.open(os.path.join(image_path, i)))
15 image = list(map(lambda i: transforms.Compose([transforms.ToTensor()])(i).
16               unsqueeze(0), image))
17
18 class Net(nn.Module):
19     def __init__(self):
20         super(Net, self).__init__()
21         self.conv1 = nn.Conv2d(3, 8, 5)
22         self.max_pool1 = nn.MaxPool2d(4, 4)
23         self.conv2 = nn.Conv2d(8, 16, 5)
24         self.max_pool2 = nn.MaxPool2d(5, 5)
25         self.conv3 = nn.Conv2d(16, 32, 5)
26         self.max_pool3 = nn.MaxPool2d(5, 5)
27         self.lin1 = nn.Linear(32, 128)
28         self.lin2 = nn.Linear(128, 84)
29         self.lin3 = nn.Linear(84, 9)
30
31     def forward(self, x):
32         x = self.conv1(x)
33         x = F.relu(x)
34         x = self.max_pool1(x)
35         x = self.conv2(x)
36         x = F.relu(x)
37         x = self.max_pool2(x)
38         x = self.conv3(x)
39         x = F.relu(x)
40         x = self.max_pool3(x)
41         x = torch.flatten(x, 1)
42         x = self.lin1(x)
```

```
41         x = F.relu(x)
42         x = self.lin2(x)
43         x = F.relu(x)
44         x = self.lin3(x)
45         return x
46 def main():
47     net = Net()
48     net.load_state_dict(torch.load('0602-657811153-Coppoletta.pt'))
49     net.eval()
50     names = ['Circle', 'Heptagon', 'Hexagon', 'Nonagon', 'Octagon', 'Pentagon', '
        Square', 'Star', 'Triangle']
51     solutions = list(zip(name, list(map(lambda k: names[k], list(map(lambda l:
        net(l).argmax().item(), image))))))
52     textfile = open("prediction.txt", "w")
53     for i in solutions:
54         print(f'{i[0]}: {i[1]}')
55         textfile.write(f'{i[0]}: {i[1]}')
56         textfile.write("\n")
57     textfile.close()
58 main()
```